

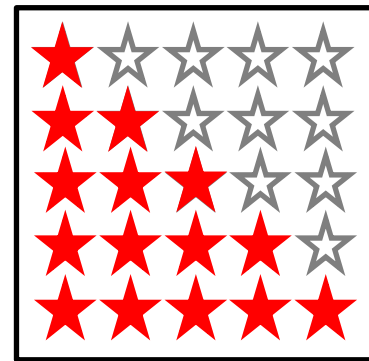
Machine Learning

Recommender Systems

Problem formulation

Example: Predicting movie ratings

User rates movies using zero to five stars



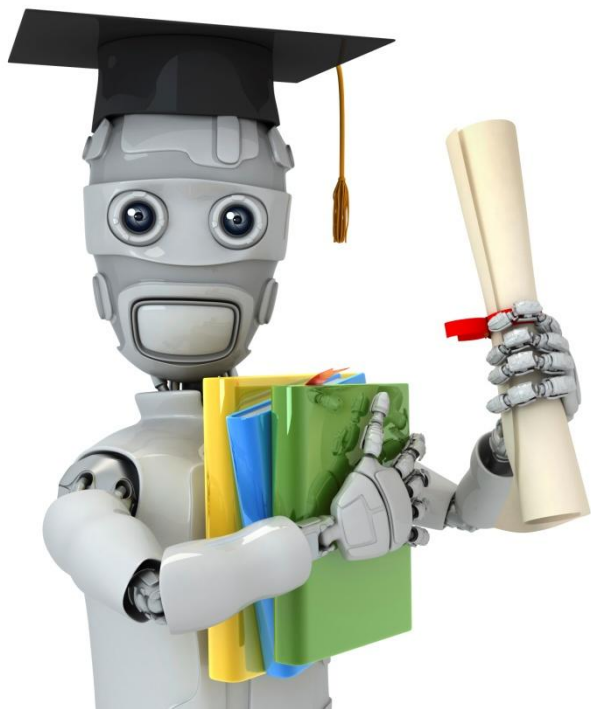
n_u = no. users

n_m = no. movies

$r(i, j) = 1$ if user j has
rated movie i

$y^{(i,j)}$ = rating given by
user j to movie i
(defined only if
 $r(i, j) = 1$)

Movie	Alice (1)	Bob (2)	Carol (3)	Dave (4)
Love at last	5	5	0	0
Romance forever	5	?	?	0
Cute puppies of love	?	4	0	?
Nonstop car chases	0	0	3	4
Swords vs. karate	0	0	5	?



Machine Learning

Recommender Systems

Content-based
recommendations

Content-based recommender systems

$n_u=4$

$n_m=5$

Movie		Alice (1) θ^1	Bob (2) θ^2	Carol (3) θ^3	Dave (4) θ^4
x^1	Love at last	5	5	0	0
x^2	Romance forever	5	?	?	0
\vdots	Cute puppies of love	?	4	0	?
	Nonstop car chases	0	0	5	4
x^5	Swords vs. karate	0	0	5	?

Let $x_0 = 1$, the intercept term, then the feature vector of movie 1, Love at last, is

$x^1 = [1, 0.9, 0]^T$, and $x^2 = [1, 1.0, 0.01]^T$, $x^5 = [1, 0, 0.9]^T$

For each user j , learn a parameter $\theta^{(j)} \in \mathbb{R}^3$. Predict user j as rating movie i with $(\theta^{(j)})^T x^{(i)}$ stars.

$\theta^{(j)} \in \mathbb{R}^{n+1}$

Example: $\theta^1 = [0, 5, 0]^T$, $x^3 = [1, 0.99, 0]^T$, then, $(\theta^1)^T x^3 = 4.95$.

Problem formulation

$r(i, j) = 1$ if user j has rated movie i (0 otherwise)

$y^{(i,j)}$ = rating by user j on movie i (if defined)

$\theta^{(j)}$ = parameter vector for user j

$x^{(i)}$ = feature vector for movie i

For user j , movie i , predicted rating: $(\theta^{(j)})^T (x^{(i)})$

$m^{(j)}$ = no. of movies rated by user j

To learn $\theta^{(j)}$:

Optimization objective:

To learn $\theta^{(j)}$ (parameter for user j):

$$\min_{\theta^{(j)}} \frac{1}{2} \sum_{i:r(i,j)=1} \left((\theta^{(j)})^T x^{(i)} - y^{(i,j)} \right)^2 + \frac{\lambda}{2} \sum_{k=1}^n (\theta_k^{(j)})^2$$

n : number of feature

To learn $\theta^{(1)}, \theta^{(2)}, \dots, \theta^{(n_u)}$:

$$\min_{\theta^{(1)}, \dots, \theta^{(n_u)}} \frac{1}{2} \sum_{j=1}^{n_u} \sum_{i:r(i,j)=1} \left((\theta^{(j)})^T x^{(i)} - y^{(i,j)} \right)^2 + \frac{\lambda}{2} \sum_{j=1}^{n_u} \sum_{k=1}^n (\theta_k^{(j)})^2$$

n_u : number of user

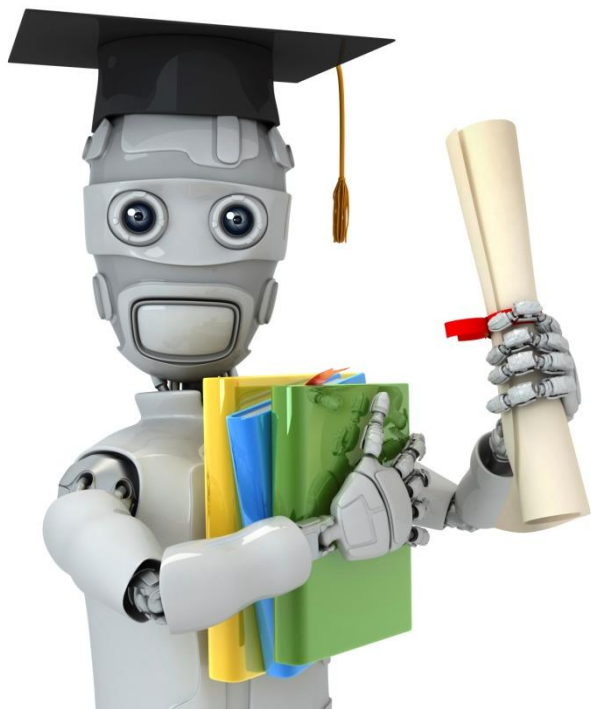
Optimization algorithm:

$$\min_{\theta^{(1)}, \dots, \theta^{(n_u)}} \underbrace{\frac{1}{2} \sum_{j=1}^{n_u} \sum_{i:r(i,j)=1} \left((\theta^{(j)})^T x^{(i)} - y^{(i,j)} \right)^2 + \frac{\lambda}{2} \sum_{j=1}^{n_u} \sum_{k=1}^n (\theta_k^{(j)})^2}_{J(\theta^{(1)}, \dots, \theta^{(n_u)})}$$

Gradient descent update:

$$\theta_k^{(j)} := \theta_k^{(j)} - \alpha \sum_{i:r(i,j)=1} ((\theta^{(j)})^T x^{(i)} - y^{(i,j)}) x_k^{(i)} \quad (\text{for } k = 0)$$

$$\theta_k^{(j)} := \theta_k^{(j)} - \alpha \underbrace{\left(\sum_{i:r(i,j)=1} ((\theta^{(j)})^T x^{(i)} - y^{(i,j)}) x_k^{(i)} + \lambda \theta_k^{(j)} \right)}_{\frac{\partial}{\partial \theta_k^{(j)}} J(\theta^{(1)}, \dots, \theta^{(n_u)})} \quad (\text{for } k \neq 0)$$



Machine Learning

Recommender Systems

Collaborative filtering

Problem motivation

Movie	Alice (1)	Bob (2)	Carol (3)	Dave (4)	x_1 (romance)	x_2 (action)
Love at last	5	5	0	0	0.9	0
Romance forever	5	?	?	0	1.0	0.01
Cute puppies of love	?	4	0	?	0.99	0
Nonstop car chases	0	0	5	4	0.1	1.0
Swords vs. karate	0	0	5	?	0	0.9

Problem motivation

Movie		Alice (1) θ^1	Bob (2) θ^2	Carol (3) θ^3	Dave (4) θ^4	x_1 (romance)	x_2 (action)
x^1	Love at last	5	5	0	0	?	?
x^2	Romance forever	5	?	?	0	?	?
\vdots	Cute puppies of love	?	4	0	?	?	?
\vdots	Nonstop car chases	0	0	5	4	?	?
x^5	Swords vs. karate	0	0	5	?	?	?

$$\theta^{(1)} = \begin{bmatrix} 0 \\ 5 \\ 0 \end{bmatrix}, \theta^{(2)} = \begin{bmatrix} 0 \\ 5 \\ 0 \end{bmatrix}, \theta^{(3)} = \begin{bmatrix} 0 \\ 0 \\ 5 \end{bmatrix}, \theta^{(4)} = \begin{bmatrix} 0 \\ 0 \\ 5 \end{bmatrix}$$

Given $\theta^{(j)}$, what feature vector should $x^{(i)}$ be?

Optimization algorithm

Given $\theta^{(1)}, \dots, \theta^{(n_u)}$, to learn $x^{(i)}$:

$$\min_{x^{(i)}} \frac{1}{2} \sum_{j:r(i,j)=1} ((\theta^{(j)})^T x^{(i)} - y^{(i,j)})^2 + \frac{\lambda}{2} \sum_{k=1}^n (x_k^{(i)})^2$$

Given $\theta^{(1)}, \dots, \theta^{(n_u)}$, to learn $x^{(1)}, \dots, x^{(n_m)}$:

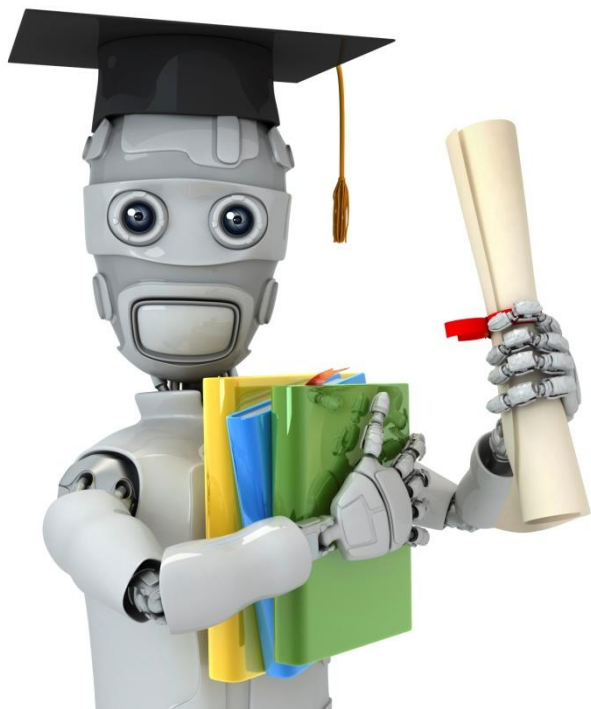
$$\min_{x^{(1)}, \dots, x^{(n_m)}} \frac{1}{2} \sum_{i=1}^{n_m} \sum_{j:r(i,j)=1} ((\theta^{(j)})^T x^{(i)} - y^{(i,j)})^2 + \frac{\lambda}{2} \sum_{i=1}^{n_m} \sum_{k=1}^n (x_k^{(i)})^2$$

Collaborative filtering

Given $x^{(1)}, \dots, x^{(n_m)}$ (and movie ratings),
can estimate $\theta^{(1)}, \dots, \theta^{(n_u)}$

Given $\theta^{(1)}, \dots, \theta^{(n_u)}$,
can estimate $x^{(1)}, \dots, x^{(n_m)}$

Randomly Guess θ , use θ learn $x(\theta \rightarrow x)$, then use x learn $\theta(x \rightarrow \theta)$.
 $\theta \rightarrow x \rightarrow \theta \rightarrow x \rightarrow \theta \rightarrow x \rightarrow \theta \rightarrow x \dots\dots\dots$



Machine Learning

Recommender Systems

Collaborative
filtering algorithm

Collaborative filtering optimization objective

Given $x^{(1)}, \dots, x^{(n_m)}$, estimate $\theta^{(1)}, \dots, \theta^{(n_u)}$:

$$\min_{\theta^{(1)}, \dots, \theta^{(n_u)}} \frac{1}{2} \sum_{j=1}^{n_u} \sum_{i:r(i,j)=1} ((\theta^{(j)})^T x^{(i)} - y^{(i,j)})^2 + \frac{\lambda}{2} \sum_{j=1}^{n_u} \sum_{k=1}^n (\theta_k^{(j)})^2$$

Given $\theta^{(1)}, \dots, \theta^{(n_u)}$, estimate $x^{(1)}, \dots, x^{(n_m)}$:

$$\min_{x^{(1)}, \dots, x^{(n_m)}} \frac{1}{2} \sum_{i=1}^{n_m} \sum_{j:r(i,j)=1} ((\theta^{(j)})^T x^{(i)} - y^{(i,j)})^2 + \frac{\lambda}{2} \sum_{i=1}^{n_m} \sum_{k=1}^n (x_k^{(i)})^2$$

Minimizing $x^{(1)}, \dots, x^{(n_m)}$ and $\theta^{(1)}, \dots, \theta^{(n_u)}$ simultaneously:

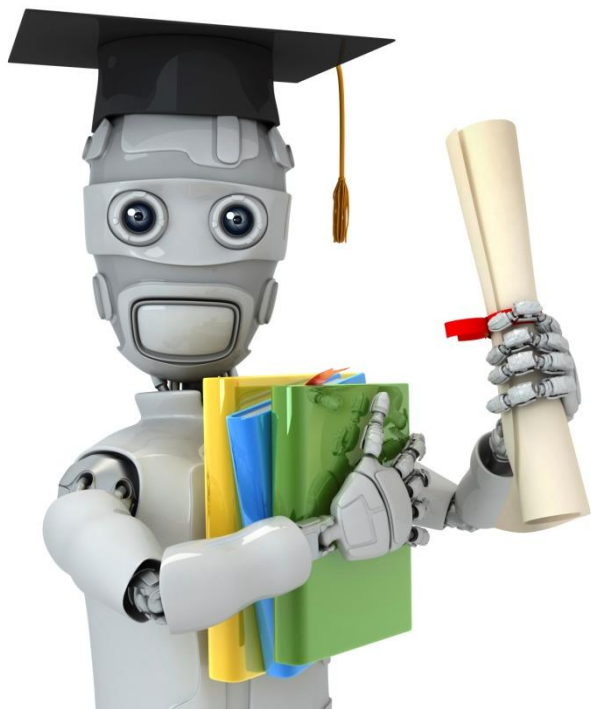
$$J(x^{(1)}, \dots, x^{(n_m)}, \theta^{(1)}, \dots, \theta^{(n_u)}) = \frac{1}{2} \sum_{(i,j):r(i,j)=1} ((\theta^{(j)})^T x^{(i)} - y^{(i,j)})^2 + \frac{\lambda}{2} \sum_{i=1}^{n_m} \sum_{k=1}^n (x_k^{(i)})^2 + \frac{\lambda}{2} \sum_{j=1}^{n_u} \sum_{k=1}^n (\theta_k^{(j)})^2$$
$$\min_{\substack{x^{(1)}, \dots, x^{(n_m)} \\ \theta^{(1)}, \dots, \theta^{(n_u)}}} J(x^{(1)}, \dots, x^{(n_m)}, \theta^{(1)}, \dots, \theta^{(n_u)})$$

Collaborative filtering algorithm

1. Initialize $x^{(1)}, \dots, x^{(n_m)}, \theta^{(1)}, \dots, \theta^{(n_u)}$ to small random values.
2. Minimize $J(x^{(1)}, \dots, x^{(n_m)}, \theta^{(1)}, \dots, \theta^{(n_u)})$ using gradient descent (or an advanced optimization algorithm). E.g. for every $j = 1, \dots, n_u, i = 1, \dots, n_m$:

$$x_k^{(i)} := x_k^{(i)} - \alpha \left(\sum_{j:r(i,j)=1} ((\theta^{(j)})^T x^{(i)} - y^{(i,j)}) \theta_k^{(j)} + \lambda x_k^{(i)} \right) \rightarrow \frac{\partial}{\partial x_k^{(i)}} J(\dots)$$
$$\theta_k^{(j)} := \theta_k^{(j)} - \alpha \left(\sum_{i:r(i,j)=1} ((\theta^{(j)})^T x^{(i)} - y^{(i,j)}) x_k^{(i)} + \lambda \theta_k^{(j)} \right) \rightarrow \frac{\partial}{\partial \theta_k^{(j)}} J(\dots)$$

3. For a user with parameters θ and a movie with (learned) features x , predict a star rating of $\theta^T x$.



Machine Learning

Recommender Systems

Vectorization:
Low rank matrix
factorization

Collaborative filtering

Movie	Alice (1)	Bob (2)	Carol (3)	Dave (4)
Love at last	5	5	0	0
Romance forever	5	?	?	0
Cute puppies of love	?	4	0	?
Nonstop car chases	0	0	5	4
Swords vs. karate	0	0	5	?

$$n_u=4 \quad n_m=5$$

$$Y = \begin{bmatrix} 5 & 5 & 0 & 0 \\ 5 & ? & ? & 0 \\ ? & 4 & 0 & ? \\ 0 & 0 & 5 & 4 \\ 0 & 0 & 5 & 0 \end{bmatrix}$$

Collaborative filtering

$$Y = \begin{bmatrix} 5 & 5 & 0 & 0 \\ 5 & ? & ? & 0 \\ ? & 4 & 0 & ? \\ 0 & 0 & 5 & 4 \\ 0 & 0 & 5 & 0 \end{bmatrix}$$

Predicted ratings:

$$\underbrace{\begin{bmatrix} \frac{(\theta^{(1)})^T(x^{(1)})}{(\theta^{(1)})^T(x^{(2)})} & (\theta^{(2)})^T(x^{(1)}) & \dots & (\theta^{(n_u)})^T(x^{(1)}) \\ (\theta^{(1)})^T(x^{(2)}) & \frac{(\theta^{(2)})^T(x^{(2)})}{(\theta^{(2)})^T(x^{(2)})} & \dots & (\theta^{(n_u)})^T(x^{(2)}) \\ \vdots & \vdots & \vdots & \vdots \\ (\theta^{(1)})^T(x^{(n_m)}) & (\theta^{(2)})^T(x^{(n_m)}) & \dots & (\theta^{(n_u)})^T(x^{(n_m)}) \end{bmatrix}}_{X\Theta^T}$$

Where

$$X = \begin{bmatrix} (x^{(1)})^T \\ (x^{(2)})^T \\ \vdots \\ (x^{(n_m)})^T \end{bmatrix}$$

$$\Theta = \begin{bmatrix} (\theta^{(1)})^T \\ (\theta^{(2)})^T \\ \vdots \\ (\theta^{(n_u)})^T \end{bmatrix}$$

Finding related movies

For each product i , we learn a feature vector $x^{(i)} \in \mathbb{R}^n$.

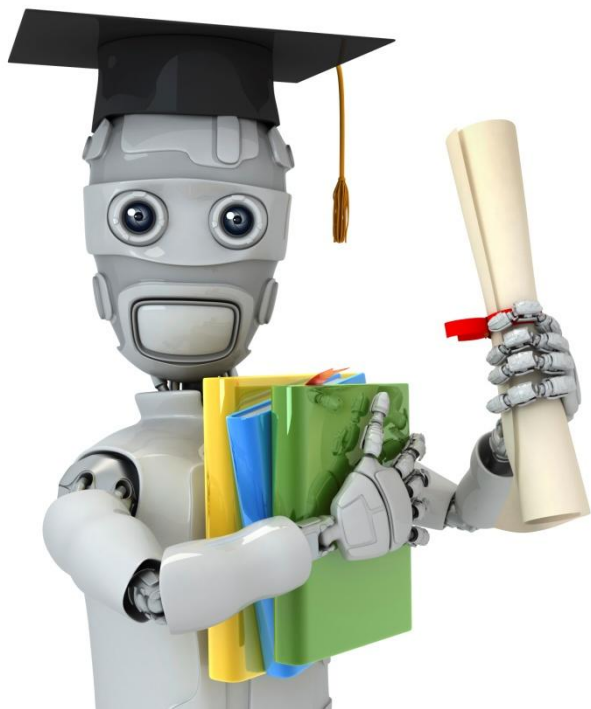
$x_1=\text{romance}, x_2=\text{action}, \dots,$

How to find movies j related to movie i ?

small $\|x^{(i)} - x^{(j)}\|$

5 most similar movies to movie i :

Find the 5 movies j with the smallest $\|x^{(i)} - x^{(j)}\|$.



Machine Learning

Recommender Systems

Implementational
detail: Mean
normalization

Users who have not rated any movies

Movie	Alice (1)	Bob (2)	Carol (3)	Dave (4)	Eve (5)
Love at last	5	5	0	0	? 0
Romance forever	5	?	?	0	? 0
Cute puppies of love	?	4	0	?	? 0
Nonstop car chases	0	0	5	4	? 0
Swords vs. karate	0	0	5	?	? 0

$$Y = \begin{bmatrix} 5 & 5 & 0 & 0 & ? \\ 5 & ? & ? & 0 & ? \\ ? & 4 & 0 & ? & ? \\ 0 & 0 & 5 & 4 & ? \\ 0 & 0 & 5 & 0 & ? \end{bmatrix}$$

$$\min_{\substack{x^{(1)}, \dots, x^{(n_m)} \\ \theta^{(1)}, \dots, \theta^{(n_u)}}} \frac{1}{2} \sum_{(i,j): r(i,j)=1} ((\theta^{(j)})^T x^{(i)} - \underline{y^{(i,j)}})^2 + \frac{\lambda}{2} \sum_{i=1}^{n_m} \sum_{k=1}^n (x_k^{(i)})^2 + \frac{\lambda}{2} \sum_{j=1}^{n_u} \sum_{k=1}^n (\theta_k^{(j)})^2$$

Let $n = 2, \theta^{(5)} \in \mathbb{R}^2$

Because no movie that Eve is rated. so $\theta^{(5)} = 0$, then for every movie, $\theta^{(5)} x^{(i)} = 0$, Eve will rate 0.

Mean Normalization:

$$Y = \begin{bmatrix} 5 & 5 & 0 & 0 & ? \\ 5 & ? & ? & 0 & ? \\ ? & 4 & 0 & ? & ? \\ 0 & 0 & 5 & 4 & ? \\ 0 & 0 & 5 & 0 & ? \end{bmatrix} \quad \mu = \begin{bmatrix} 2.5 \\ 2.5 \\ 2 \\ 2.25 \\ 1.25 \end{bmatrix} \rightarrow Y = \begin{bmatrix} 2.5 & 2.5 & -2.5 & -2.5 & ? \\ 2.5 & ? & ? & -2.5 & ? \\ ? & 2 & -2 & ? & ? \\ -2.25 & -2.25 & 2.75 & 1.75 & ? \\ -1.25 & -1.25 & 3.75 & -1.25 & ? \end{bmatrix}$$

For user j , on movie i predict:

$$(\theta^{(j)})^T (x^{(i)}) + u_i$$

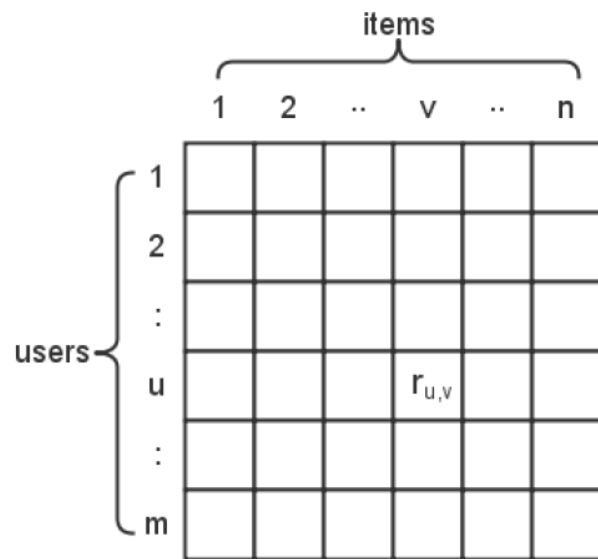
User 5 (Eve):

$$\theta^{(5)} = \begin{bmatrix} 0 \\ 0 \end{bmatrix}$$

$$(\theta^{(j)})^T (x^{(i)}) + u_i = u_i$$

Matrix Factorization

- Matrix Factorization is an effective method for recommender systems.



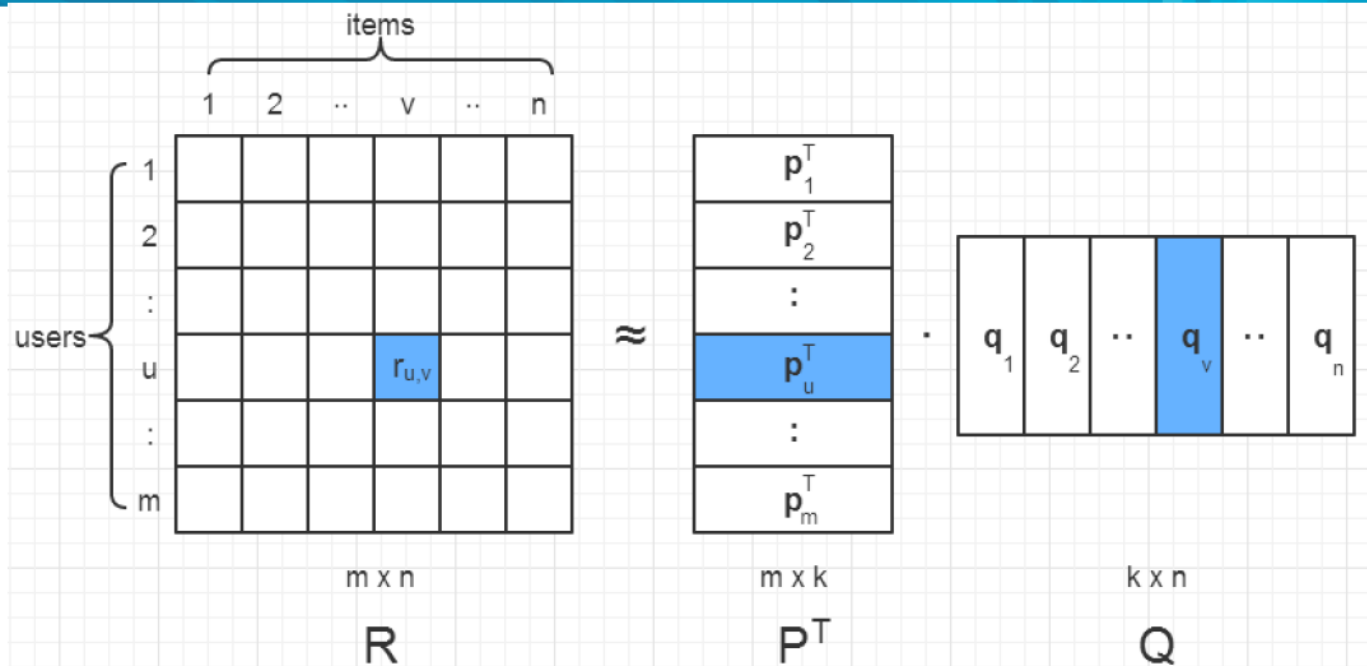
m, n : numbers of users and items

u, v : index for u_{th} user and v_{th} item

$r_{u,v}$: u_{th} user gives a rating $r_{u,v}$ to v_{th} item

libmf1[1]-2-310 (2).png $m \times n$ (7)

Matrix Factorization (Cont'd)



k : number of latent dimensions

$$r_{u,v} = \mathbf{p}_u^T \mathbf{q}_v$$

$$\text{Objective function: } \min_{P, Q} \sum_{(u,v) \in R} (r_{u,v} - \mathbf{p}_u^T \mathbf{q}_v)^2 + \lambda_P \|\mathbf{p}_u\|_F^2 + \lambda_Q \|\mathbf{q}_v\|_F^2 \quad (1)$$

Stochastic Gradient Descent (SGD)

- The basic idea of SGD is that, instead of expensively calculating the gradient of

$$\min_{P,Q} \sum_{(u,v) \in R} (r_{u,v} - \mathbf{p}_u^T \mathbf{q}_v)^2 + \lambda_P \|\mathbf{p}_u\|_F^2 + \lambda_Q \|\mathbf{q}_v\|_F^2$$

, it randomly selects a (u,v) entry from the summation and calculates the corresponding gradient.

- Once $r_{u,v}$ is chosen, the objective function is reduced to

$$(r_{u,v} - \mathbf{p}_u^T \mathbf{q}_v)^2 + \lambda_P \mathbf{p}_u^T \mathbf{p}_u + \lambda_Q \mathbf{q}_v^T \mathbf{q}_v$$

SGD (Cont'd)

$$J = (r_{u,v} - \mathbf{p}_u^T \mathbf{q}_v)^2 + \lambda_P \mathbf{p}_u^T \mathbf{p}_u + \lambda_Q \mathbf{q}_v^T \mathbf{q}_v$$

$$J = (e_{u,v})^2 + \lambda_P \mathbf{p}_u^T \mathbf{p}_u + \lambda_Q \mathbf{q}_v^T \mathbf{q}_v \quad \text{where } e_{u,v} = (r_{u,v} - \mathbf{p}_u^T \mathbf{q}_v)$$

$$\frac{\partial J}{\partial \mathbf{p}_u} = 2e_{u,v}(-\mathbf{q}_v) + 2\lambda_P \mathbf{p}_u = -2(e_{u,v} \mathbf{q}_v - \lambda_P \mathbf{p}_u)$$

$$\frac{\partial J}{\partial \mathbf{q}_v} = 2e_{u,v}(-\mathbf{p}_u) + 2\lambda_Q \mathbf{q}_v = -2(e_{u,v} \mathbf{p}_u - \lambda_Q \mathbf{q}_v)$$

$$\mathbf{p}_u = \mathbf{p}_u + \gamma(e_{u,v} \mathbf{q}_v - \lambda_P \mathbf{p}_u)$$

$$\mathbf{q}_v = \mathbf{q}_v + \gamma(e_{u,v} \mathbf{p}_u - \lambda_Q \mathbf{q}_v)$$

SGD (Cont'd)

$$J = (r_{u,v} - \mathbf{p}_u^T \mathbf{q}_v)^2 + \lambda_P \mathbf{p}_u^T \mathbf{p}_u + \lambda_Q \mathbf{q}_v^T \mathbf{q}_v$$

$$J = (e_{u,v})^2 + \lambda_P \mathbf{p}_u^T \mathbf{p}_u + \lambda_Q \mathbf{q}_v^T \mathbf{q}_v \quad \text{where } e_{u,v} = (r_{u,v} - \mathbf{p}_u^T \mathbf{q}_v)$$

consider the user bias and item bias :

ub : the user bias, **ub_u** : the bias of u_{th} user

ib : the item bias, **ib_i** : the bias of i_{th} item

avg : the average of all rates

so $e_{u,v}$ now equals $r_{u,v} - (\mathbf{p}_u^T \mathbf{q}_v + avg + \mathbf{ub}_u + \mathbf{ib}_v)$

SGD (Cont'd)

$$J = [r_{u,v} - (\mathbf{p}_u^T \mathbf{q}_v + \text{avg} + \mathbf{ub}_u + \mathbf{ib}_v)]^2 + \lambda_P \mathbf{p}_u^T \mathbf{p}_u + \lambda_Q \mathbf{q}_v^T \mathbf{q}_v + \lambda_{ub} \mathbf{ub}_u^2 + \lambda_{ib} \mathbf{ib}_v^2$$

$$J = (e_{u,v})^2 + \lambda_P \mathbf{p}_u^T \mathbf{p}_u + \lambda_Q \mathbf{q}_v^T \mathbf{q}_v + \lambda_P \mathbf{ub}_u^2 + \lambda_Q \mathbf{ib}_v^2$$

$$\text{where } e_{u,v} = r_{u,v} - (\mathbf{p}_u^T \mathbf{q}_v + \text{avg} + \mathbf{ub}_u + \mathbf{ib}_v)$$

$$\frac{\partial J}{\partial \mathbf{p}_u} = 2e_{u,v}(-\mathbf{q}_v) + 2\lambda_P \mathbf{p}_u = -2(e_{u,v} \mathbf{q}_v - \lambda_P \mathbf{p}_u), \quad \mathbf{p}_u = \mathbf{p}_u + \gamma(e_{u,v} \mathbf{q}_v - \lambda_P \mathbf{p}_u)$$

$$\frac{\partial J}{\partial \mathbf{q}_v} = 2e_{u,v}(-\mathbf{p}_u) + 2\lambda_Q \mathbf{q}_v = -2(e_{u,v} \mathbf{p}_u - \lambda_Q \mathbf{q}_v), \quad \mathbf{q}_v = \mathbf{q}_v + \gamma(e_{u,v} \mathbf{p}_u - \lambda_Q \mathbf{q}_v)$$

$$\frac{\partial J}{\partial \mathbf{ub}_u} = 2e_{u,v}(-1) + 2\lambda_{ub} \mathbf{ub}_u = -2(e_{u,v} - \lambda_{ub} \mathbf{ub}_u), \quad \mathbf{ub}_u = \mathbf{ub}_u + \gamma(e_{u,v} - \lambda_{ub} \mathbf{ub}_u)$$

$$\frac{\partial J}{\partial \mathbf{ib}_v} = 2e_{u,v}(-1) + 2\lambda_{ib} \mathbf{ib}_v = -2(e_{u,v} - \lambda_{ib} \mathbf{ib}_v), \quad \mathbf{ib}_v = \mathbf{ib}_v + \gamma(e_{u,v} - \lambda_{ib} \mathbf{ib}_v)$$